

MATH 680 Homework 3 Fall 2015

This homework is due on Monday, November 9 at 11:59pm. Submit your solutions (for problems 1--3) in a pdf document on MyCourses.

Please also submit text files with your R code (for problem 1--3), which must be commented and properly indented.

For problem 4, submit the R package file which should have the suffix `.tar.gz`.

1. Please see homework 1 for definitions of $X, y, \tilde{X}, \tilde{y}$ and β_* . Suppose we want to compute the bridge penalized least-squares estimate of β_* defined by

$$\hat{\beta}_{-1}^{(\lambda, \alpha)} \in \arg \min_{\tilde{\beta} \in \mathbb{R}^{p-1}} \left\{ \frac{1}{2} \|\tilde{y} - \tilde{X} \tilde{\beta}\|^2 + \frac{\lambda}{\alpha} \sum_{j=1}^{p-1} |\tilde{\beta}_j|^\alpha \right\} \quad (1)$$

$$\hat{\beta}_1^{(\lambda, \alpha)} = \bar{y} - \bar{x}' \hat{\beta}_{-1}^{(\lambda, \alpha)},$$

where $\lambda \geq 0$ and $\alpha \in [1, 2]$ are tuning parameters. When $\alpha = 2$, (1) becomes ridge-penalized least-squares. When $\alpha = 1$, (1) becomes lasso-penalized least squares.

- (a) Let f be the objective function in (1). Is f convex for $\alpha \in [1, 2]$? Explain.
- (b) Suppose that $\lambda > 0$. Find the set $S_1 \subset [1, 2]$ for which f is differentiable when $\alpha \in S_1$. Find the set $S_2 \subset [1, 2]$ for which f is twice differentiable when $\alpha \in S_2$. Provide justification.
- (c) Create an algorithm to compute $\hat{\beta}^{(\lambda, \alpha)} = (\hat{\beta}_1^{(\lambda, \alpha)}, \hat{\beta}_{-1}^{(\lambda, \alpha)})'$ for $\alpha \in (1, 2)$. Write an R function to test this algorithm. This function should have six arguments:
 - `X`, the design matrix in $\mathbb{R}^{n \times p}$, with ones in its first column.
 - `y`, the response vector in \mathbb{R}^n
 - `lam`, this is λ .
 - `alpha`, this is $\alpha \in (1, 2)$.
 - `tol`, this is a value used to determine algorithm convergence.
 - `maxit`, this is the maximum number of iterations allowed.

and return a list with three elements:

- `b1`, this is $\hat{\beta}_1^{(\lambda, \alpha)} \in \mathbb{R}$
 - `b`, this is $\hat{\beta}_{-1}^{(\lambda, \alpha)} \in \mathbb{R}^{p-1}$
 - `total.iterations`, this is the number of iterations the algorithm performed.
- (d) Pick reasonable values for n, p, β_*, ω_* and X and randomly generate data from the Normal linear regression cases model to test your function. Use your R function to compute $\hat{\beta}_{\lambda, \alpha}$ for some positive value of λ and some $\alpha \in (1, 2)$. Is the final iterate approximately a stationary point of f ?

2. In Binomial logistic regression, we assume that the observed response y_i for the i th case is the observed value of the sample proportion of successes in n_i independent and identical Bernoulli trials where the explanatory variables had values $x_i = (1, x_{i2}, \dots, x_{ip})'$. Specifically, y_i is a realization of Y_i where

$$n_i Y_i \sim \text{Binom} \left(n_i, \frac{\exp(x_i' \beta_*)}{1 + \exp(x_i' \beta_*)} \right), \quad i = 1, \dots, n;$$

Y_1, \dots, Y_n are independent; and $\beta_* \in \mathbb{R}^p$ is the unknown vector of regression coefficients. Then

$$\begin{aligned} P(Y_i = y_i) &= P(n_i Y_i = n_i y_i) \\ &= \binom{n_i}{n_i y_i} \left[\frac{\exp(x'_i \beta_*)}{1 + \exp(x'_i \beta_*)} \right]^{n_i y_i} \left[1 - \frac{\exp(x'_i \beta_*)}{1 + \exp(x'_i \beta_*)} \right]^{n_i - n_i y_i} \end{aligned}$$

where $y_i = 0/n_i, 1/n_i, \dots, n_i/n_i$. The non-random log-likelihood function $\ell : \mathbb{R}^p \rightarrow \mathbb{R}$ is defined by

$$\begin{aligned} \ell(\beta) &= \sum_{i=1}^n n_i y_i \log \left[\frac{\exp(x'_i \beta)}{1 + \exp(x'_i \beta)} \right] \\ &\quad + \sum_{i=1}^n (n_i - n_i y_i) \log \left[1 - \frac{\exp(x'_i \beta)}{1 + \exp(x'_i \beta)} \right] + \sum_{i=1}^n \log \binom{n_i}{n_i y_i}. \end{aligned}$$

Define the Bridge-penalized logistic regression estimate of β_* by

$$\bar{\beta}^{(\lambda, \alpha)} = \arg \min_{\beta \in \mathbb{R}^p} \left\{ -\ell(\beta) + \frac{\lambda}{\alpha} \sum_{j=2}^p |\beta_j|^\alpha \right\}, \quad (2)$$

where $\lambda \geq 0$ and $\alpha \in [1, 2]$ are tuning parameters.

- (a) Let g be the objective function in (2). Is g convex for $\alpha \in [1, 2]$? Explain.
- (b) Create an algorithm to compute $\bar{\beta}_\lambda$ for $\alpha \in (1, 2)$. Write an R function to test this algorithm. This function should have six arguments:
 - **X**, the model matrix in $\mathbb{R}^{n \times p}$, with ones in its first column.
 - **y**, the response vector in \mathbb{R}^n , the observed sample proportions.
 - **n.list**, a vector of n entries with i th entry n_i .
 - **lam**, this is λ .
 - **alpha**, this is $\alpha \in (1, 2)$.
 - **tol**, this is a value used to determine algorithm convergence.
 - **maxit**, this is the maximum number of iterations allowed.

and return a list with two elements:

- **b**, this is the final iterate for $\bar{\beta}^{(\lambda, \alpha)}$.
 - **total.iterations**, this is the number of iterations the algorithm performed.
- (c) Pick reasonable values for $n, n_1, \dots, n_n, p, x_1, \dots, x_n$, and β_* and generate data from the Binomial logistic regression model. Use your R function to compute $\bar{\beta}^{(\lambda, \alpha)}$ for some positive value of λ and some value $\alpha \in (1, 2)$. Is this final iterate approximately a stationary point of g ?

3. In this problem you will fit a Gaussian mixture model and use it to cluster n observations x_1, \dots, x_n , where $x_i \in \mathbb{R}^p$. The model assumes that x_1, \dots, x_n are a realization of the sequence of independent random vectors X_1, \dots, X_n defined in the following way. For each X_i there

is an unobservable random variable Y_i and $(X_1, Y_1), \dots, (X_n, Y_n)$ are independent copies of (X, Y) where Y is supported on $\{1, \dots, C\}$ and $(X|Y = y) \sim N_p(\mu_{*y}, \Sigma_{*y})$ for $y \in \{1, \dots, C\}$. Let $P(Y = y) = \pi_{*y}$ for $y \in \{1, \dots, C\}$. To fit this model, we estimate

$$\theta_* = (\pi_{*1}, \dots, \pi_{*C}, \mu_{*1}, \dots, \mu_{*C}, \Sigma_{*1}, \dots, \Sigma_{*C}).$$

Let $f_X(\cdot; \theta)$ be the density for X with respect to ν_X (lebesgue measure on \mathbb{R}^p), let $f_Y(\cdot; \theta)$ be the density for Y with respect to ν_Y (counting measure on $\{1, \dots, C\}$), and let $f_{X|Y}(\cdot; y, \theta)$ be the density of X given $Y = y$ with respect to ν_X , i.e. $f_{X|Y}(\cdot; y, \theta) = \phi(\cdot; \mu_y, \Sigma_y)$ is the density for $N_p(\mu_y, \Sigma_y)$. We have that

$$f_X(x; \theta) = \int f_{X|Y}(x; y, \theta) f_Y(y; \theta) \nu_Y(dy) = \sum_{y=1}^C \phi(x, \mu_y, \Sigma_y) \pi_y.$$

So to estimate θ_* with a realization x_1, \dots, x_n of n independent copies of X , we maximize the loglikelihood function, i.e. we compute a local minimizer for

$$\arg \min_{\theta \in \Theta} - \left[\sum_{i=1}^n \log \left\{ \sum_{y=1}^C \phi(x_i; \mu_y, \Sigma_y) \pi_y \right\} \right], \quad (3)$$

where $\theta = (\pi_1, \dots, \pi_C, \mu_1, \dots, \mu_C, \Sigma_1, \dots, \Sigma_C)$ and $\Theta = \{r \in (0, 1)^C : \sum_{j=1}^p r_j = 1\} \times \mathbb{R}^p \times \dots \times \mathbb{R}^p \times \mathbb{S}_+^p \times \dots \times \mathbb{S}_+^p$.

- (a) Write an R function that uses the EM algorithm to compute a local minimizer for (3). Choose appropriate arguments and return values for the function.
- (b) Derive an EM algorithm to compute a local minimizer for

$$\arg \min_{\theta \in \Theta} - \left[\sum_{i=1}^n \log \left\{ \sum_{y=1}^C \phi(x_i; \mu_y, \Sigma) \pi_y \right\} \right], \quad (4)$$

where $\theta = (\pi_1, \dots, \pi_C, \mu_1, \dots, \mu_C, \Sigma)$ and $\Theta = \{r \in (0, 1)^C : \sum_{j=1}^p r_j = 1\} \times \mathbb{R}^p \times \dots \times \mathbb{R}^p \times \mathbb{S}_+^p$. The objective function in (4) is the negative loglikelihood function for the model where $(X|Y = y) \sim N_p(\mu_{*y}, \Sigma_*)$, so the covariance matrix is the same for all $y \in \{1, \dots, C\}$. Write the corresponding R function to compute this local minimizer.

- (c) Find a classification dataset from the UCI machine learning data repository for which the Gaussian mixture model assumptions are not terribly unreasonable (pretending that the responses/class labels are unobserved). This repository is at

<https://archive.ics.uci.edu/ml/datasets.html>

Test the R functions created in parts a and b on these data by pretending that the responses/class labels y_1, \dots, y_n are unobserved. Set C equal to the number possible categories for a response/class label. For a fitted Gaussian mixture model, assign a class label \hat{y}_i to the i th observation x_i with

$$\hat{y}_i = \arg \max_{y \in \{1, \dots, C\}} \hat{P}(Y = y | X = x_i),$$

where $\hat{P}(Y = y|X = x_i)$ is an estimate of $P(Y = y|X = x_i)$ from the Gaussian mixture model fit. Using the fit from the R function created in part a, report $n^{-1} \sum_{i=1}^n 1(\hat{y}_i = y_i)$. Also do this for the fit from the R function created in part b.

4. Create an R package called `mylasso` that allows the user to call two functions: `minlassoLSwithC` and `minlassoLS`, where these functions are defined in “rlasso.r” posted on moodle.
 - The definition of `minlassoLSwithC` requires the C function defined in “rlasso.c” that is also posted on moodle.
 - This package should have three help files: one for the function `minlassoLSwithC`, one for the function `minlassoLS`, and one for the R package `mylasso`.
 - Build the package on a linux machine using R CMD `build`. Linux machines are available in 352 Ford Hall and could be accessed on your personal computer using an ssh client.
 - Check the package with R CMD `check` and make sure that it passes all of the tests (copy and paste the output).
 - Install your built package using R CMD `INSTALL` and make sure that the functions work properly and that the help files are displaying properly.
 - Upload the built package file, which should have a suffix of `.tar.gz`